TMA Innovation Center

Data Science Level 1

Hung Nguyen

TIC Data Team Lead

Course objective

Theory

Deeply understand about Data Science:

- What is DS
- What DS [can] do
- How to be a DS

Practice

- Build an end-to-end predictive model

Course syllabus

Session 1

- Data Science introduction
- Data science project cycle
- Basic concepts of statistics
- Intro duction to Python

Session 2

- Data manipulation with Python - **Pandas**
- Data Visualization with Python - **Matplotlib**

Session 3

- Machine Learning with **Scikit-Learn** – build model
- Model evaluation
- Basic algorithms
- PoCs planning

Session 4

- PoCs review:
  - ➢ Prediction with linear regression
  - ➢ Classification with logistic regression

TMA Innovation Center

Data Science Level 1

-- Session 1--

Introduction

Hung Nguyen

TIC Data Team Lead


0101
1010
0010
Data
Science

What is Data Science?

- Data science is a **multi-disciplinary** field
- Uses scientific methods, processes, algorithms and systems
- Extracts **knowledge** and **insights** from *data in various forms*, both structured and unstructured

Harvard Business Review

DATA SCIENTIST

The Sexiest Job of the 21st Century

Josh Wills
@josh_wills

Follow

Data Scientist (n.): Person who is better at statistics than any software engineer and better at software engineering than any statistician.

Reply Retweet Favorite More

9:55 AM - 3 May 12

What Data Scientist do?

Product and Research

- Find **better** techniques
- Make sure ML system was **“smart” enough**
- **Help** the **engineering** team
- Work with product managers to **incorporate new improvements**
- Develop **new** products or features

Sales and Marketing

- **Explain** technology aspects to **potential clients**
- Work **with data set** from a potential client
- **Share** knowledge
- Create **visualization**

Data Scientist skill set

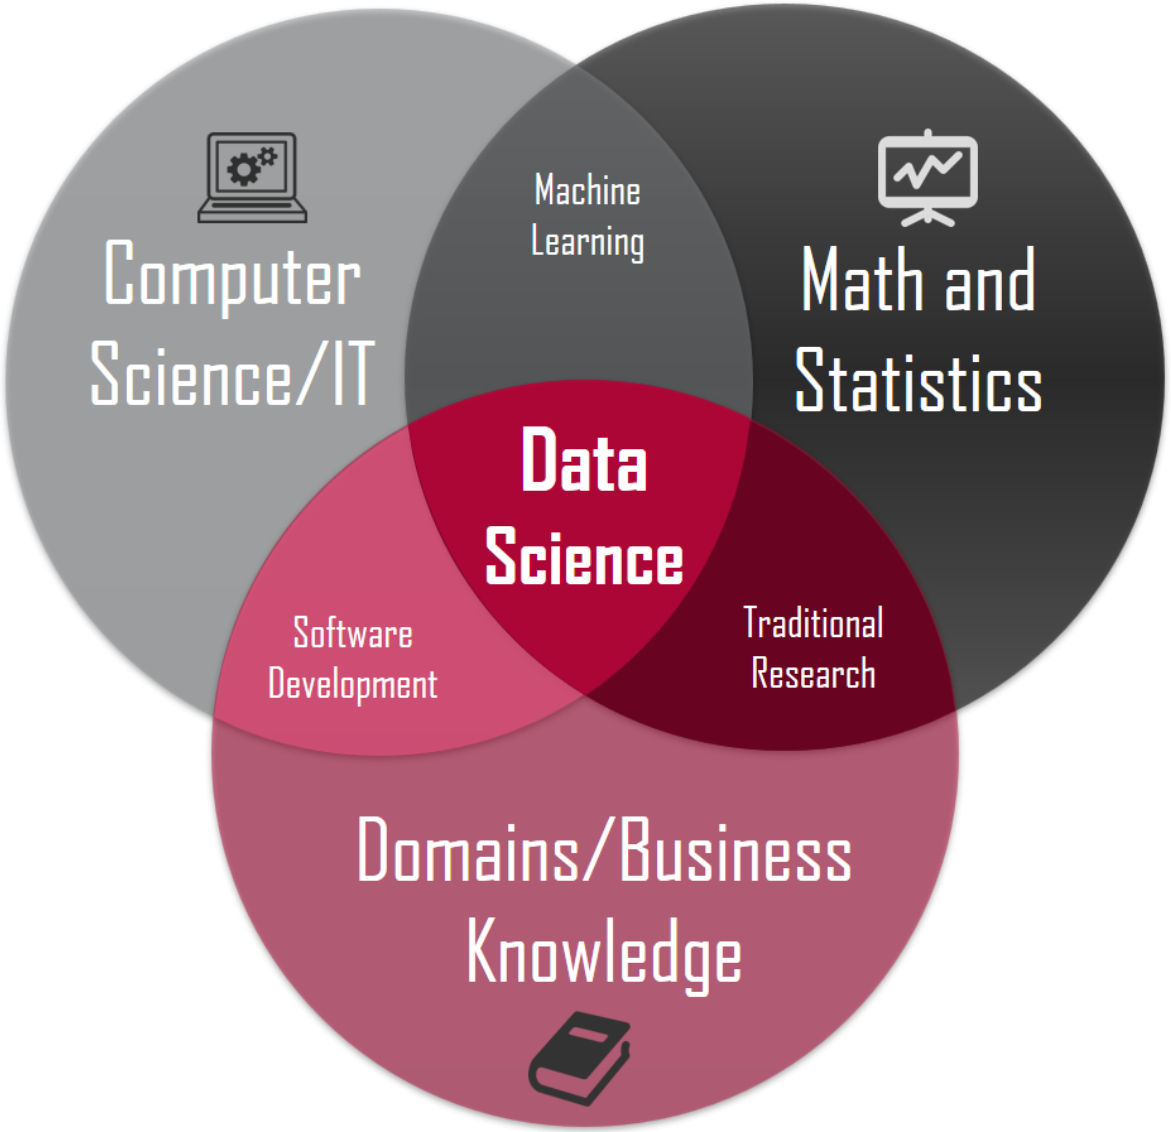

Data Scientist skill set

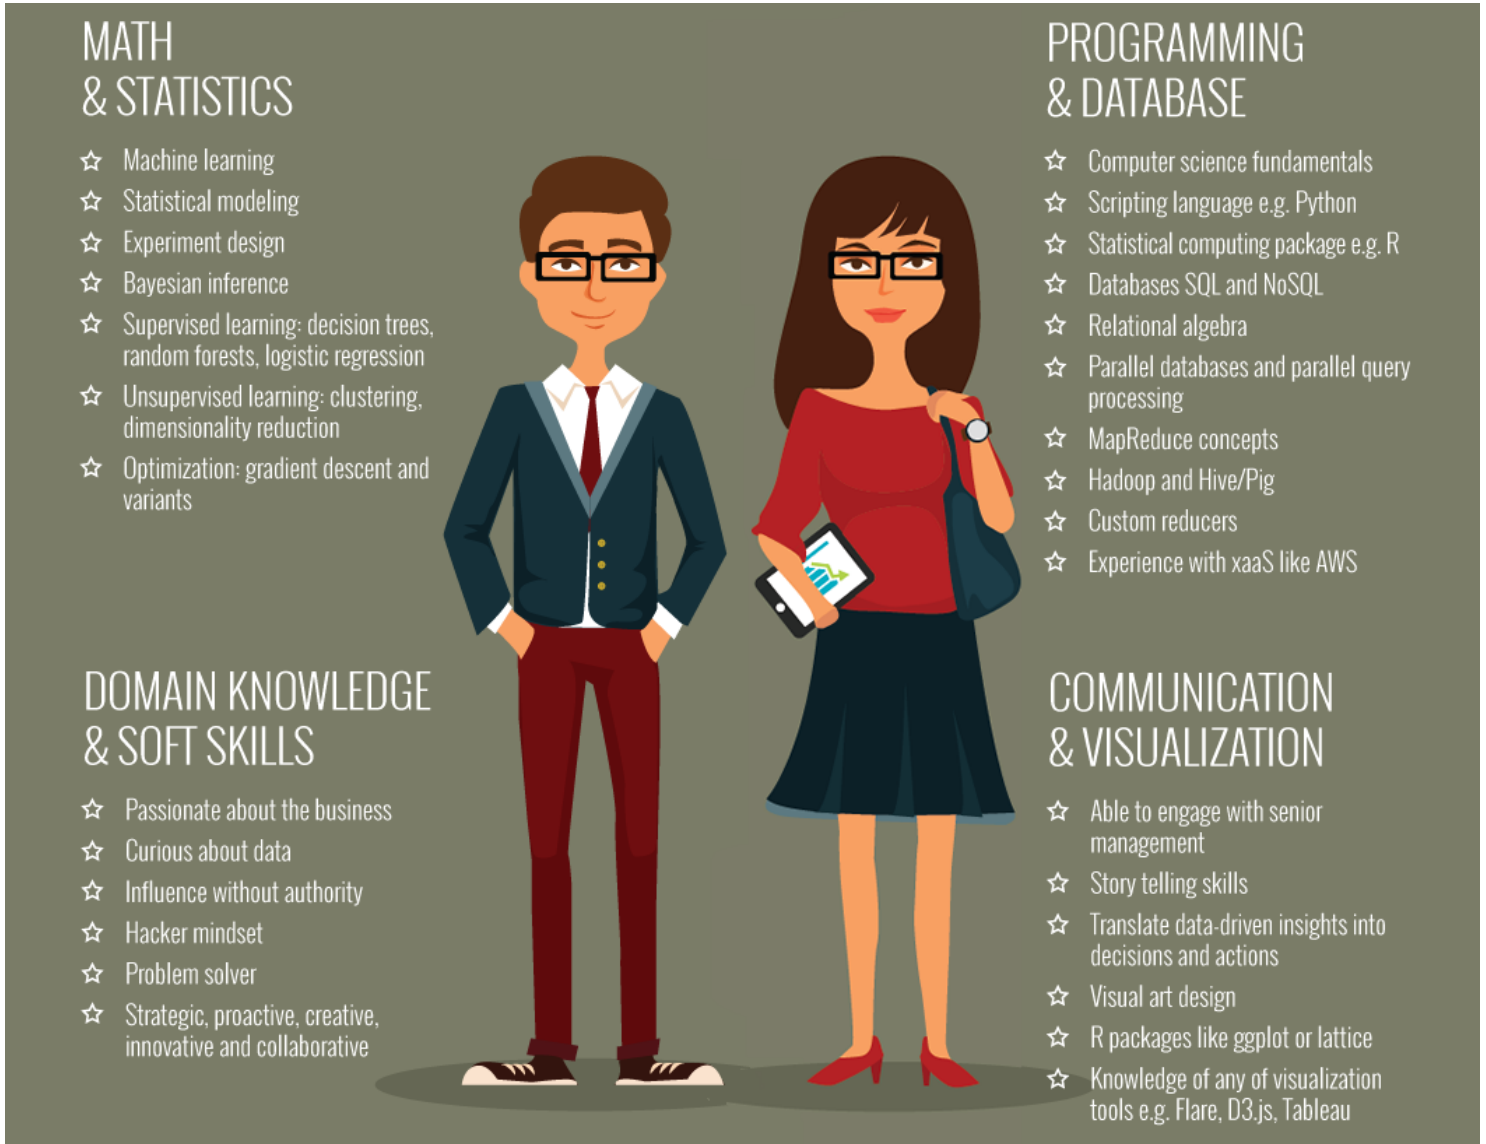

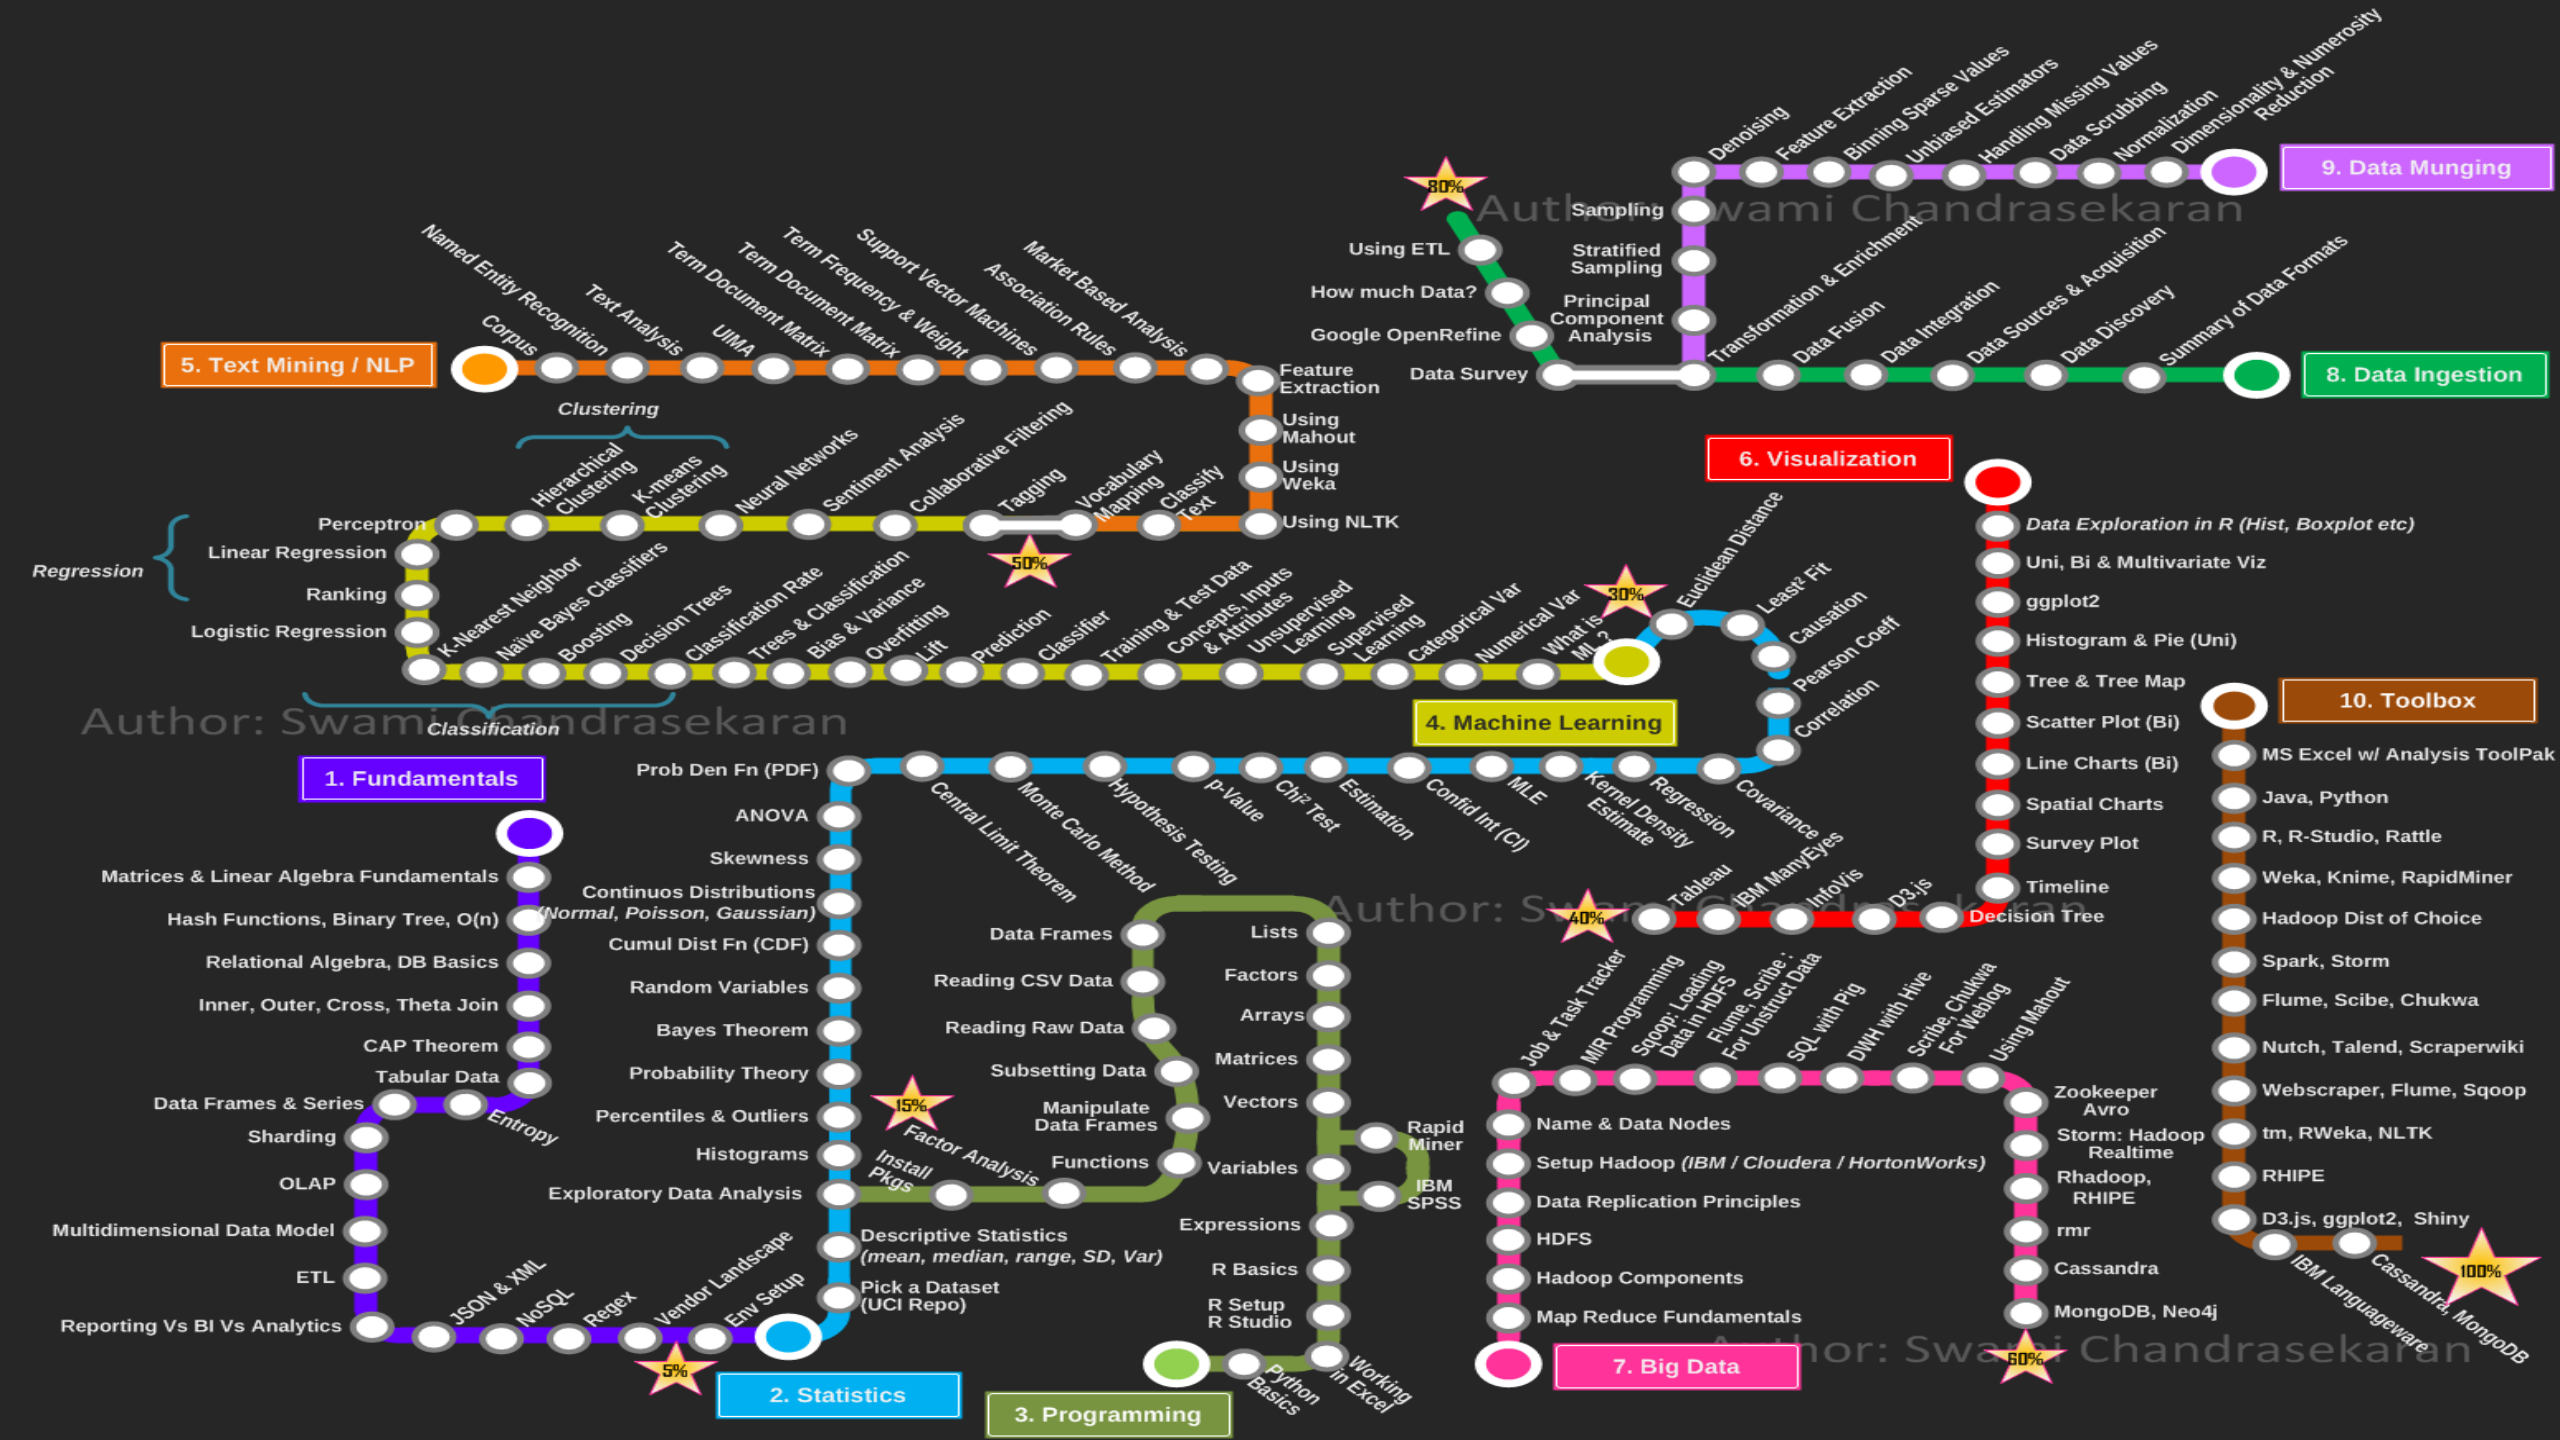
Author: Swami Chandrasekaran
1. Fundamentals
Matrices & Linear Algebra Fundamentals
Hash Functions, Binary Tree, O(n)
Relational Algebra, DB Basics
Inner, Outer, Cross, Theta Join
CAP Theorem
Tabular Data
Data Frames & Series
Entropy
Sharding
OLAP
Multidimensional Data Model
ETL
Reporting Vs BI Vs Analytics
JSON & XML
NoSQL
Regex
Vendor Landscape
Env Setup
5%
2. Statistics
Prob Den Fn (PDF)
ANOVA
Skewness
Continuos Distributions (Normal, Poisson, Gaussian)
Cumul Dist Fn (CDF)
Random Variables
Bayes Theorem
Probability Theory
Percentiles & Outliers
Histograms
Exploratory Data Analysis
Descriptive Statistics (mean, median, range, SD, Var)
Pick a Dataset (UCI Repo)
Central Limit Theorem
Monte Carlo Method
Hypothesis Testing
p-Value
Chi² Test
Estimation
Confid Int (CI)
MLE
Kernel Density Estimate
Regression
Covariance
15%
Factor Analysis
Install Pkgs
3. Programming
Data Frames
Reading CSV Data
Reading Raw Data
Subsetting Data
Manipulate Data Frames
Functions
Lists
Factors
Arrays
Matrices
Vectors
Variables
Expressions
R Basics
R Setup
R Studio
Rapid Miner
IBM SPSS
Python Basics
Working in Excel
4. Machine Learning
Regression
Perceptron
Linear Regression
Ranking
Logistic Regression
Classification
K-Nearest Neighbor
Naive Bayes Classifiers
Boosting
Decision Trees
Classification Rate
Trees & Classification
Bias & Variance
Overfitting
Lift
Prediction
Classifier
Training & Test Data
Concepts, Inputs & Attributes
Unsupervised Learning
Supervised Learning
Categorical Var
Numerical Var
What is ML ?
30%
Euclidean Distance
Least² Fit
Causation
Pearson Coeff
Correlation
Clustering
Hierarchical Clustering
K-means Clustering
Neural Networks
Sentiment Analysis
Collaborative Filtering
Tagging
Vocabulary Mapping
Classify Text
50%
5. Text Mining / NLP
Corpus
Named Entity Recognition
Text Analysis
UIMA
Term Document Matrix
Term Document Matrix
Term Frequency & Weight
Support Vector Machines
Association Rules
Market Based Analysis
Feature Extraction
Using Mahout
Using Weka
Using NLTK
6. Visualization
Data Exploration in R (Hist, Boxplot etc)
Uni, Bi & Multivariate Viz
ggplot2
Histogram & Pie (Uni)
Tree & Tree Map
Scatter Plot (Bi)
Line Charts (Bi)
Spatial Charts
Survey Plot
Timeline
Decision Tree
40%
Tableau
IBM ManyEyes
InfoVis
D3.js
7. Big Data
Job & Task Tracker
MR Programming
Sqoop: Loading Data in HDFS
Flume, Scribe : For Unstruct Data
SQL with Pig
DWH with Hive
Scribe, Chukwa For Weblog
Using Mahout
Name & Data Nodes
Setup Hadoop (IBM / Cloudera / HortonWorks)
Data Replication Principles
HDFS
Hadoop Components
Map Reduce Fundamentals
Zookeeper Avro
Storm: Hadoop Realtime
Rhadoop, RHIPE
rmr
Cassandra
MongoDB, Neo4j
60%
8. Data Ingestion
Using ETL
How much Data?
Google OpenRefine
Data Survey
Sampling
Stratified Sampling
Principal Component Analysis
80%
Transformation & Enrichment
Data Fusion
Data Integration
Data Sources & Acquisition
Data Discovery
Summary of Data Formats
9. Data Munging
Denoising
Feature Extraction
Binning Sparse Values
Unbiased Estimators
Handling Missing Values
Data Scrubbing
Normalization
Dimensionality & Numerosity Reduction
10. Toolbox
MS Excel w/ Analysis ToolPak
Java, Python
R, R-Studio, Rattle
Weka, Knime, RapidMiner
Hadoop Dist of Choice
Spark, Storm
Flume, Scibe, Chukwa
Nutch, Talend, Scraperwiki
Webscraper, Flume, Sqoop
tm, RWeka, NLTK
RHIPE
D3.js, ggplot2, Shiny
IBM Languageware
Cassandra, MongoDB
100%

TMA Innovation Center

Data Science Level 1

-- Session 1--

Data science project cycle

Hung Nguyen

TIC Data Team Lead

The Data Science Process

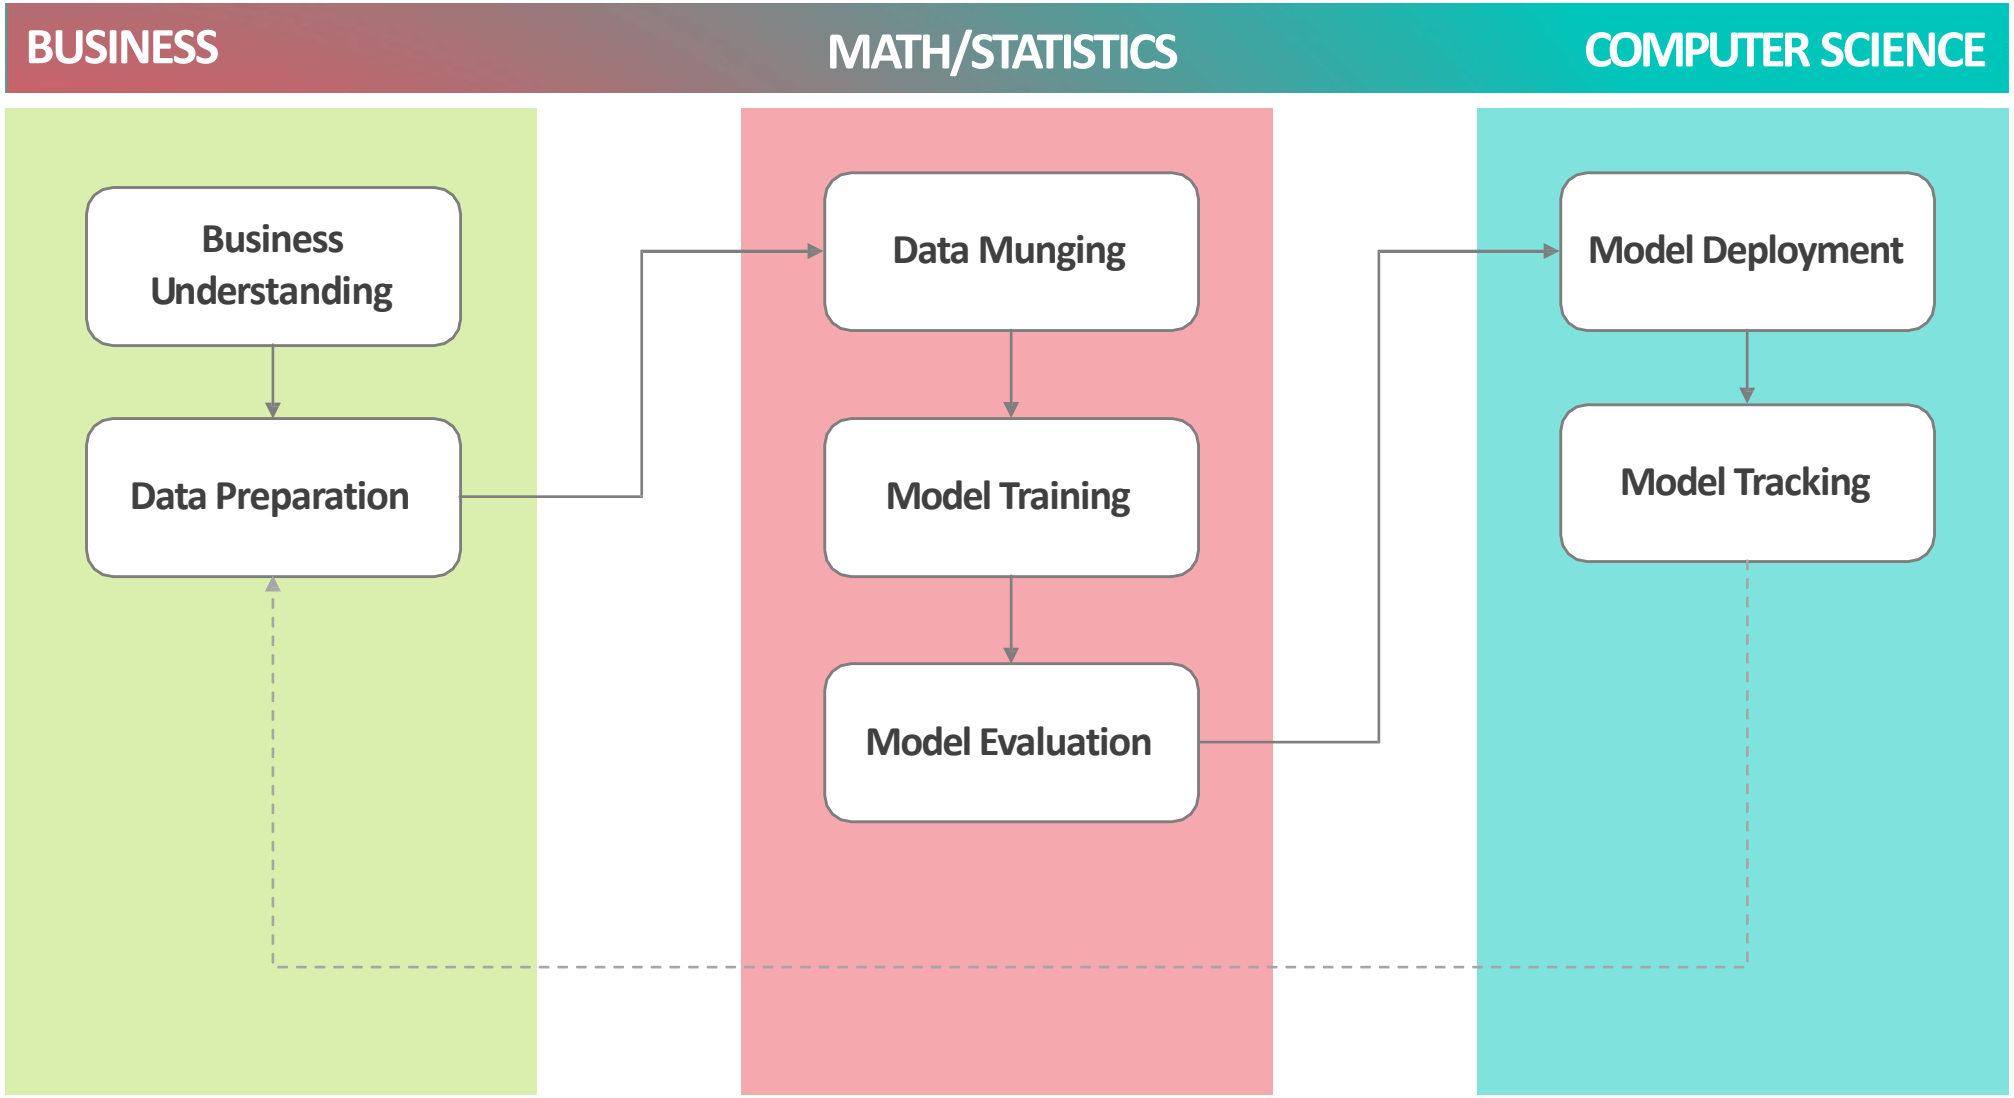

Business understanding

Determine

What does the client want to achieve?

- ✓ Reduce attrition
- ✓ Customized targeting
- ✓ Plan future media spend
- ✓ Prevent fraud
- ✓ Recommend Products

Understand

- **Understand** success criteria.
- **List** assumptions, constraints, and important factors.
- **Identify** secondary or competing objectives.
- **Study** existing solutions (if any).

Map

- **Business Objective → Technical Objective**
  - ✓ State the project objective in **technical terms**.
  - ✓ How data science project will help
  - ✓ **Successful** scenarios.

Data Preparation

IDENTIFY

- Data **sources, formats**
- Entity **Relationship** Diagram
- Identify **relevant** data
- Record **unavailable** data.
- How long

COLLECT

- Access or acquire all relevant data in **a central location**
- Quality control **checks** and **tests**

ASSESS

- Get **familiar** with the data.
- Study **seasonality**.
- Detect **mistakes**.
- Check **assumptions**.
- Review **distributions**.

VECTORIZE

- Create the Analysis **Dataset**

Data munging (preprocessing)

- **Descriptive** statistics
- **Correlation** analysis
- Impute **missing** values
- Trim **extreme** values
- Process **categorical** attributes
- **Transformations** (square, log, etc.)
- Multicollinearity: **reduce** redundancy
- Create **additional feature**
- **Interactions**
- **Normalization** (scaling)

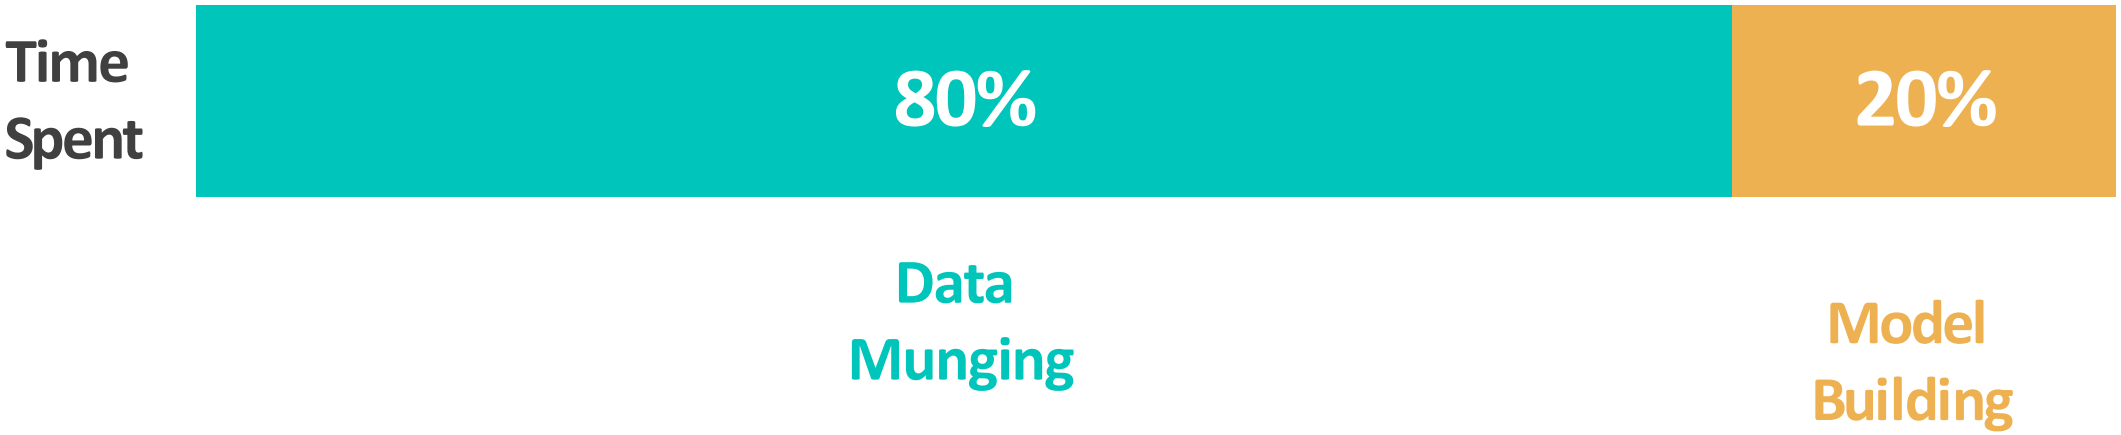
Time Spent
80%
20%
Data Munging
Model Building

Model training

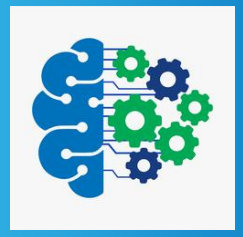

Try **more than one** machine learning technique.

Fine-tune parameters.

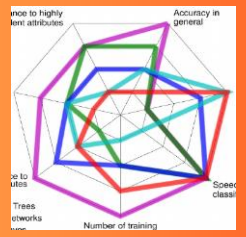

Assess model **performance**.

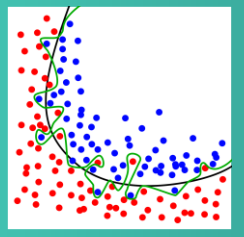

Avoid **Over-fitting**.

Model Evaluation

MODEL SELECTION

- Law of Parsimony:

simple is better

- Execution time
- Deployment complexity

ASSESSMENT

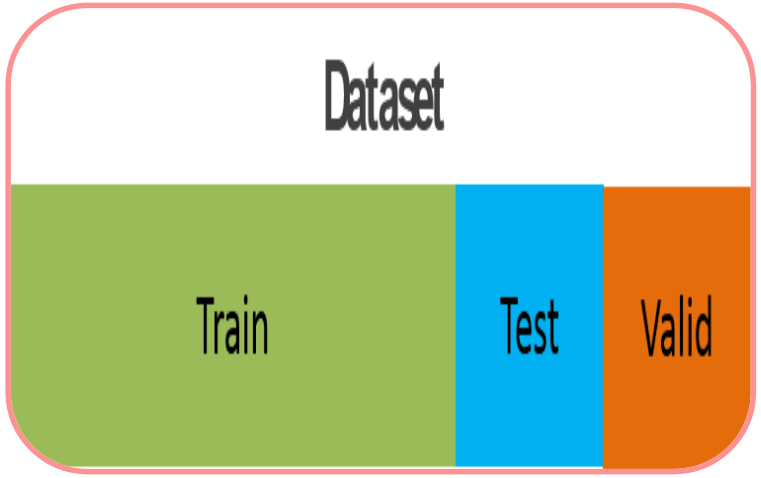

PRESENTATION

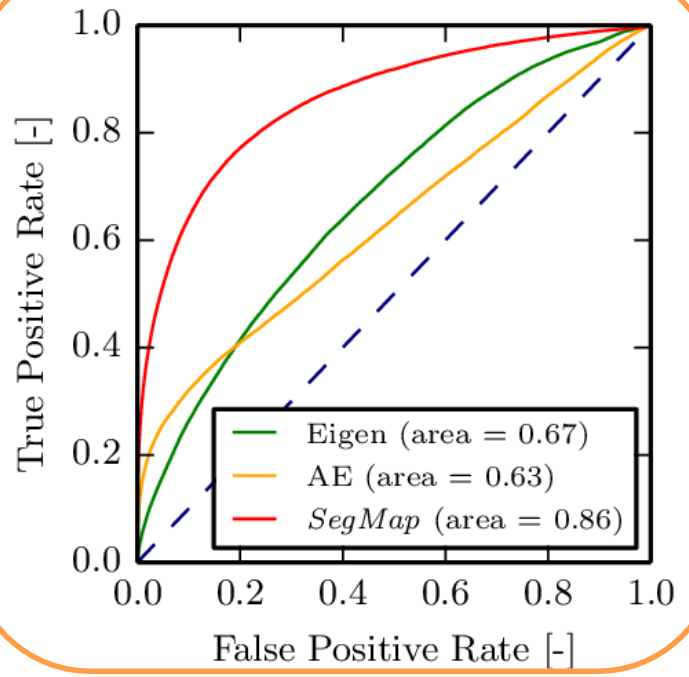

Model Deployment

- Model production cycle
- Scoring code, or publish model as a web service
- Model Documentation (Technical Specifications)
- Reproducibility
- Model Persistence vs. Model Transience

Model Tracking

MONITOR

- Model performance over time
- Predictor distribution

MAINTAIN

- Model maintenance plan
- Adding new data sources
- Version control

TEST

- Experimental Design (A/B tests, Fractional Factorial)

Data Science Process: Recap

Business Understanding	Data Preparation	Data Munging	Model Training	Model Evaluation	Model Deployment	Model Tracking
Determine	Identify	Impute	Train	Evaluate	Deploy	Monitor
Understand	Collect	Transform	Assess	Peer Review	Document	Maintain
Map	Assess	Reduce	Select	Present		Test
	Vectorize					

DISCUSS COLLATE WRANGLE PERFORM COMMUNICATE EXECUTE TRACK

TMA Innovation Center

Data Science Level 1

-- Session 1--

Basic statistical concepts

Hung Nguyen

TIC Data Team Lead

Statistical Features

- min, max, mean, median, mode, standard deviation, first & third quartiles
- the first stats technique to apply when exploring a dataset
- easy to understand and implement in code

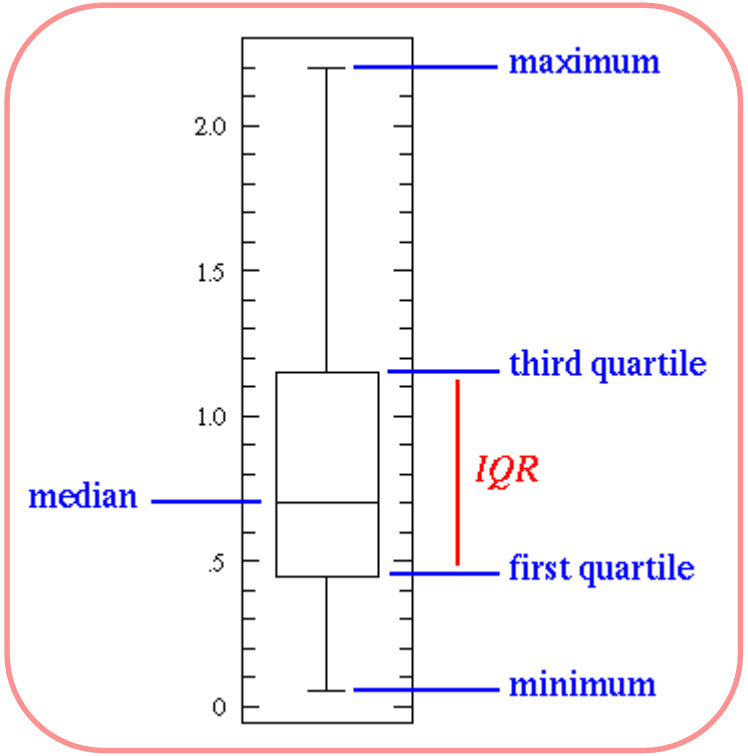

Probability Distributions

Uniform

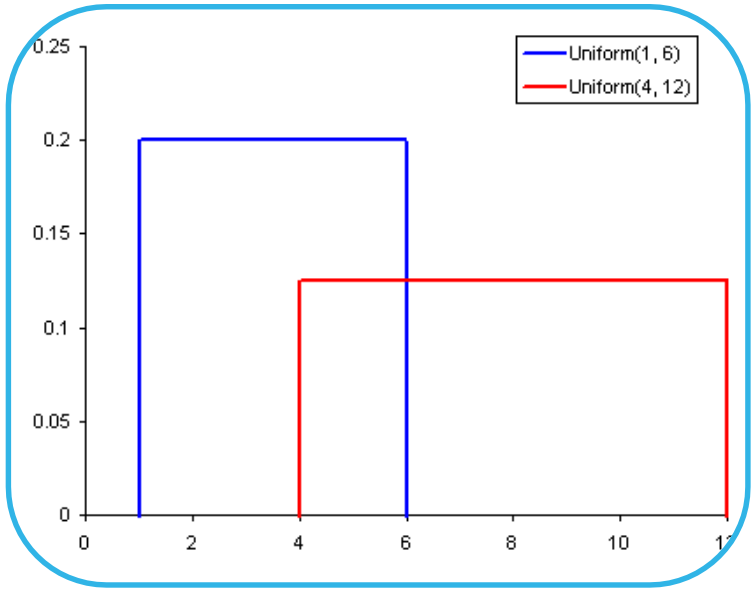

Gaussian

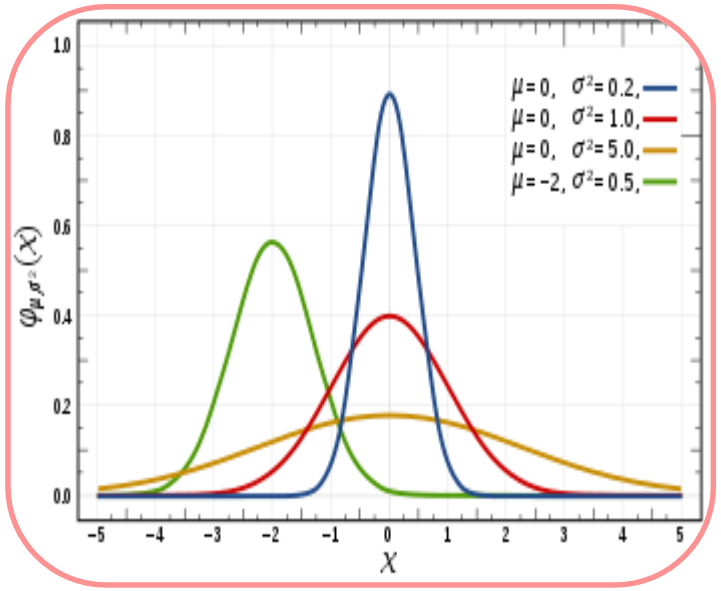

Poisson

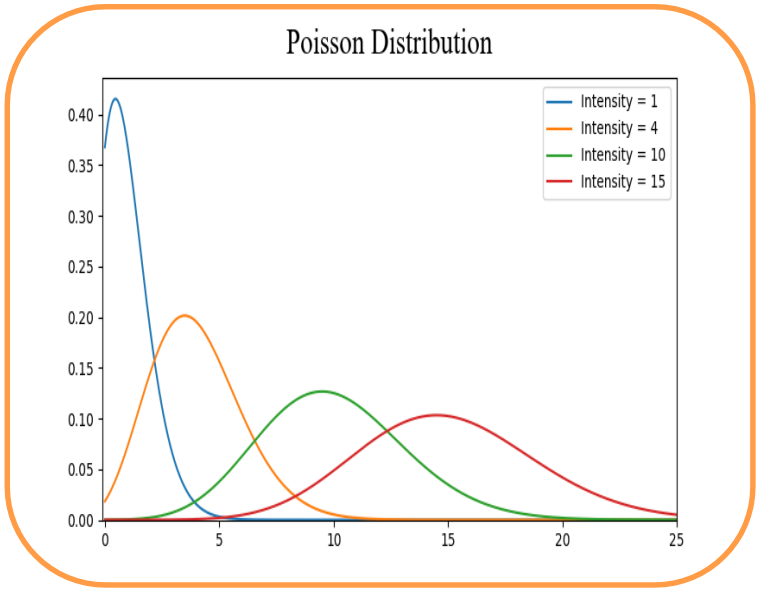

Probability - percent chance that some event will occur

Dimensionality Reduction

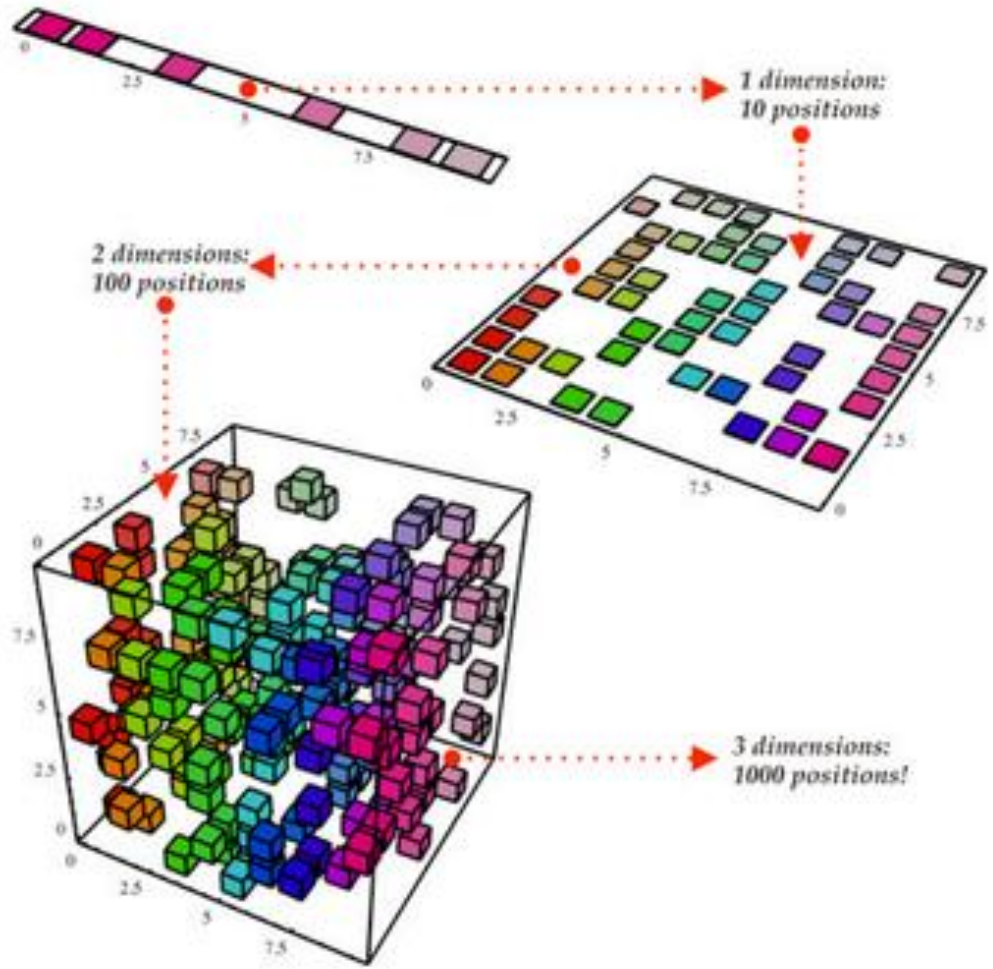

Techniques:

- Feature pruning
- Principal component analysis (PCA)

Over and Under Sampling

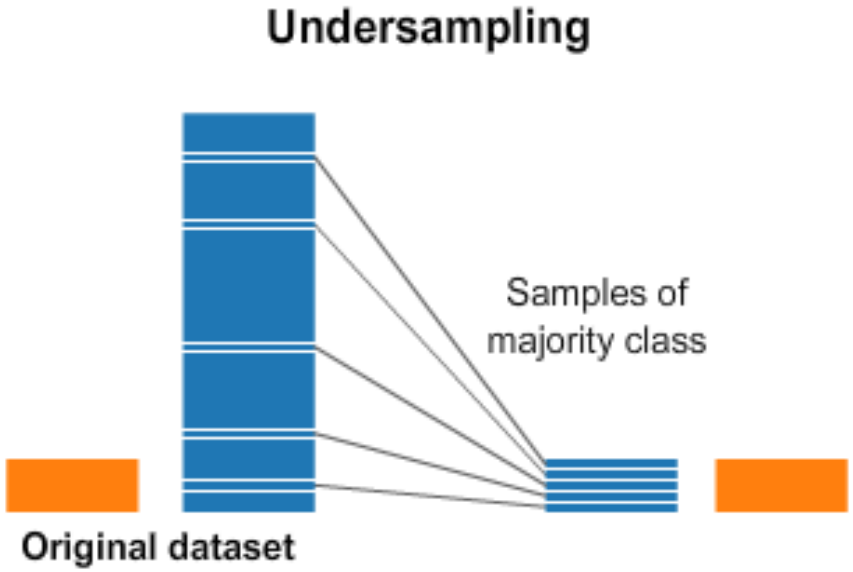

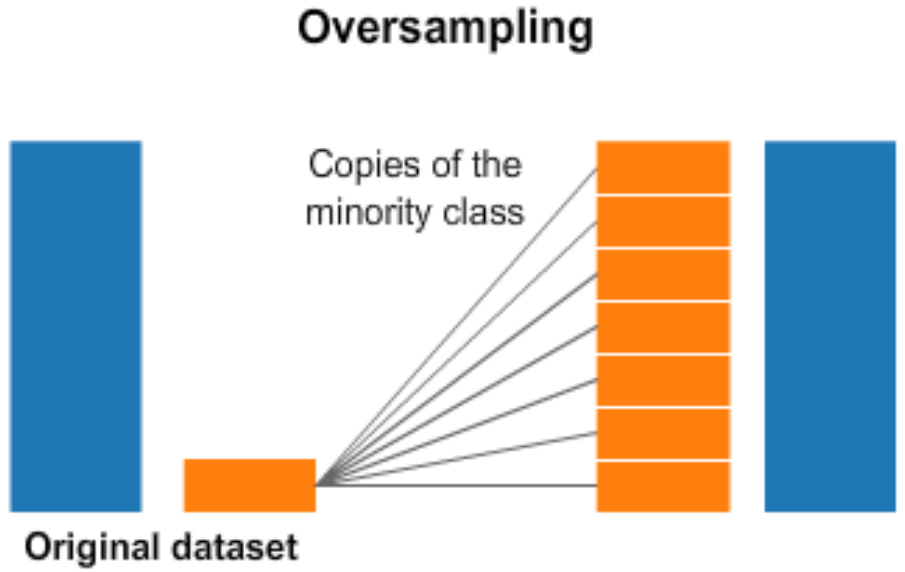

usually use in **imbalanced classification** problem

Bayesian Statistics

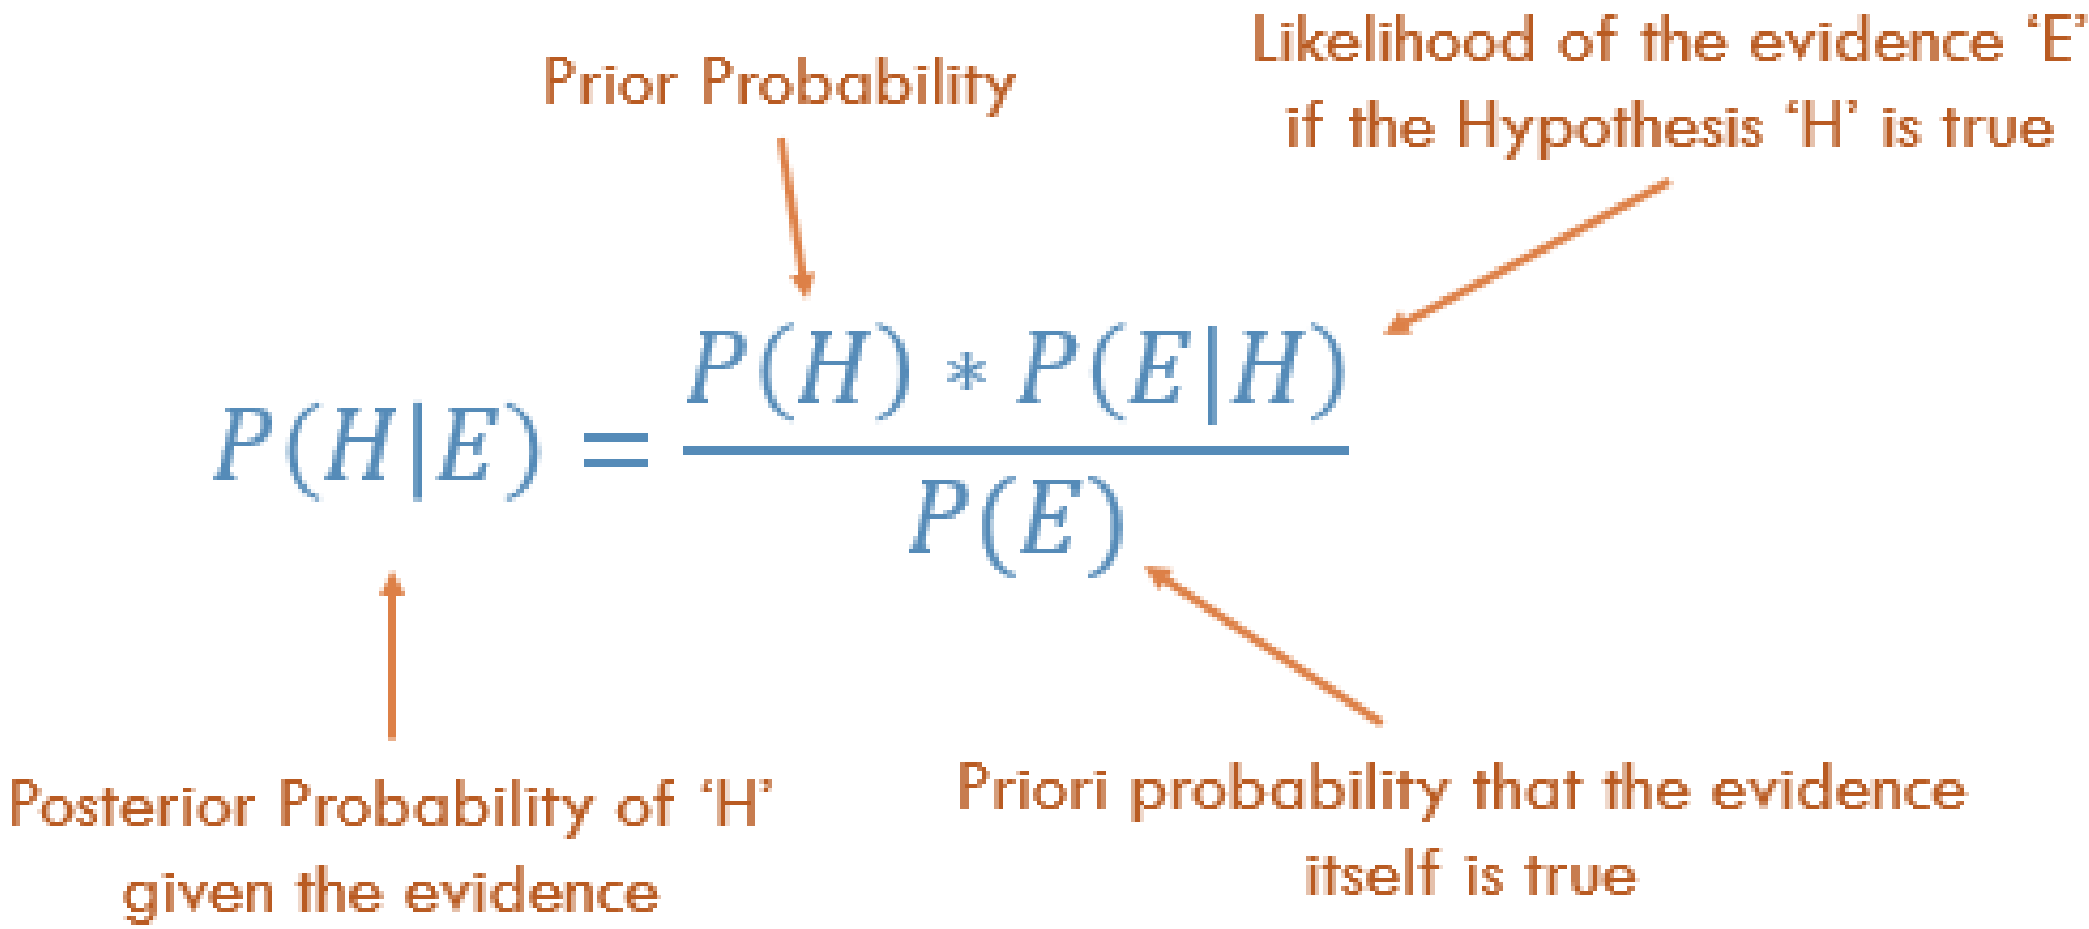

- Example 1 : the probability of a certain medical test being positive is 90%, if a patient has disease D. 1% of the population have the disease, and the test records a false positive 5% of the time. If you receive a positive test, what is your probability of having D?

- P(+|D)=0.9, P(D)=0.01, P(+|no D)=0.05, we want P(D|+)

- $$P(D|+) = \frac{P(+|D)\,P(D)}{P(+)} = \frac{P(+|D)P(D)}{P(+|D)P(D) + P(+|\mathrm{no}D)P(\mathrm{no}D)}$$

- Substituting in the numbers : P(D|+) = 0.15

TMA Innovation Center

Data Science Level 1

-- Session 1--

Introduction to Python

Hung Nguyen

TIC Data Team Lead

Python, R, Both, or Other platforms for Analytics, Data Science, Machine Learning

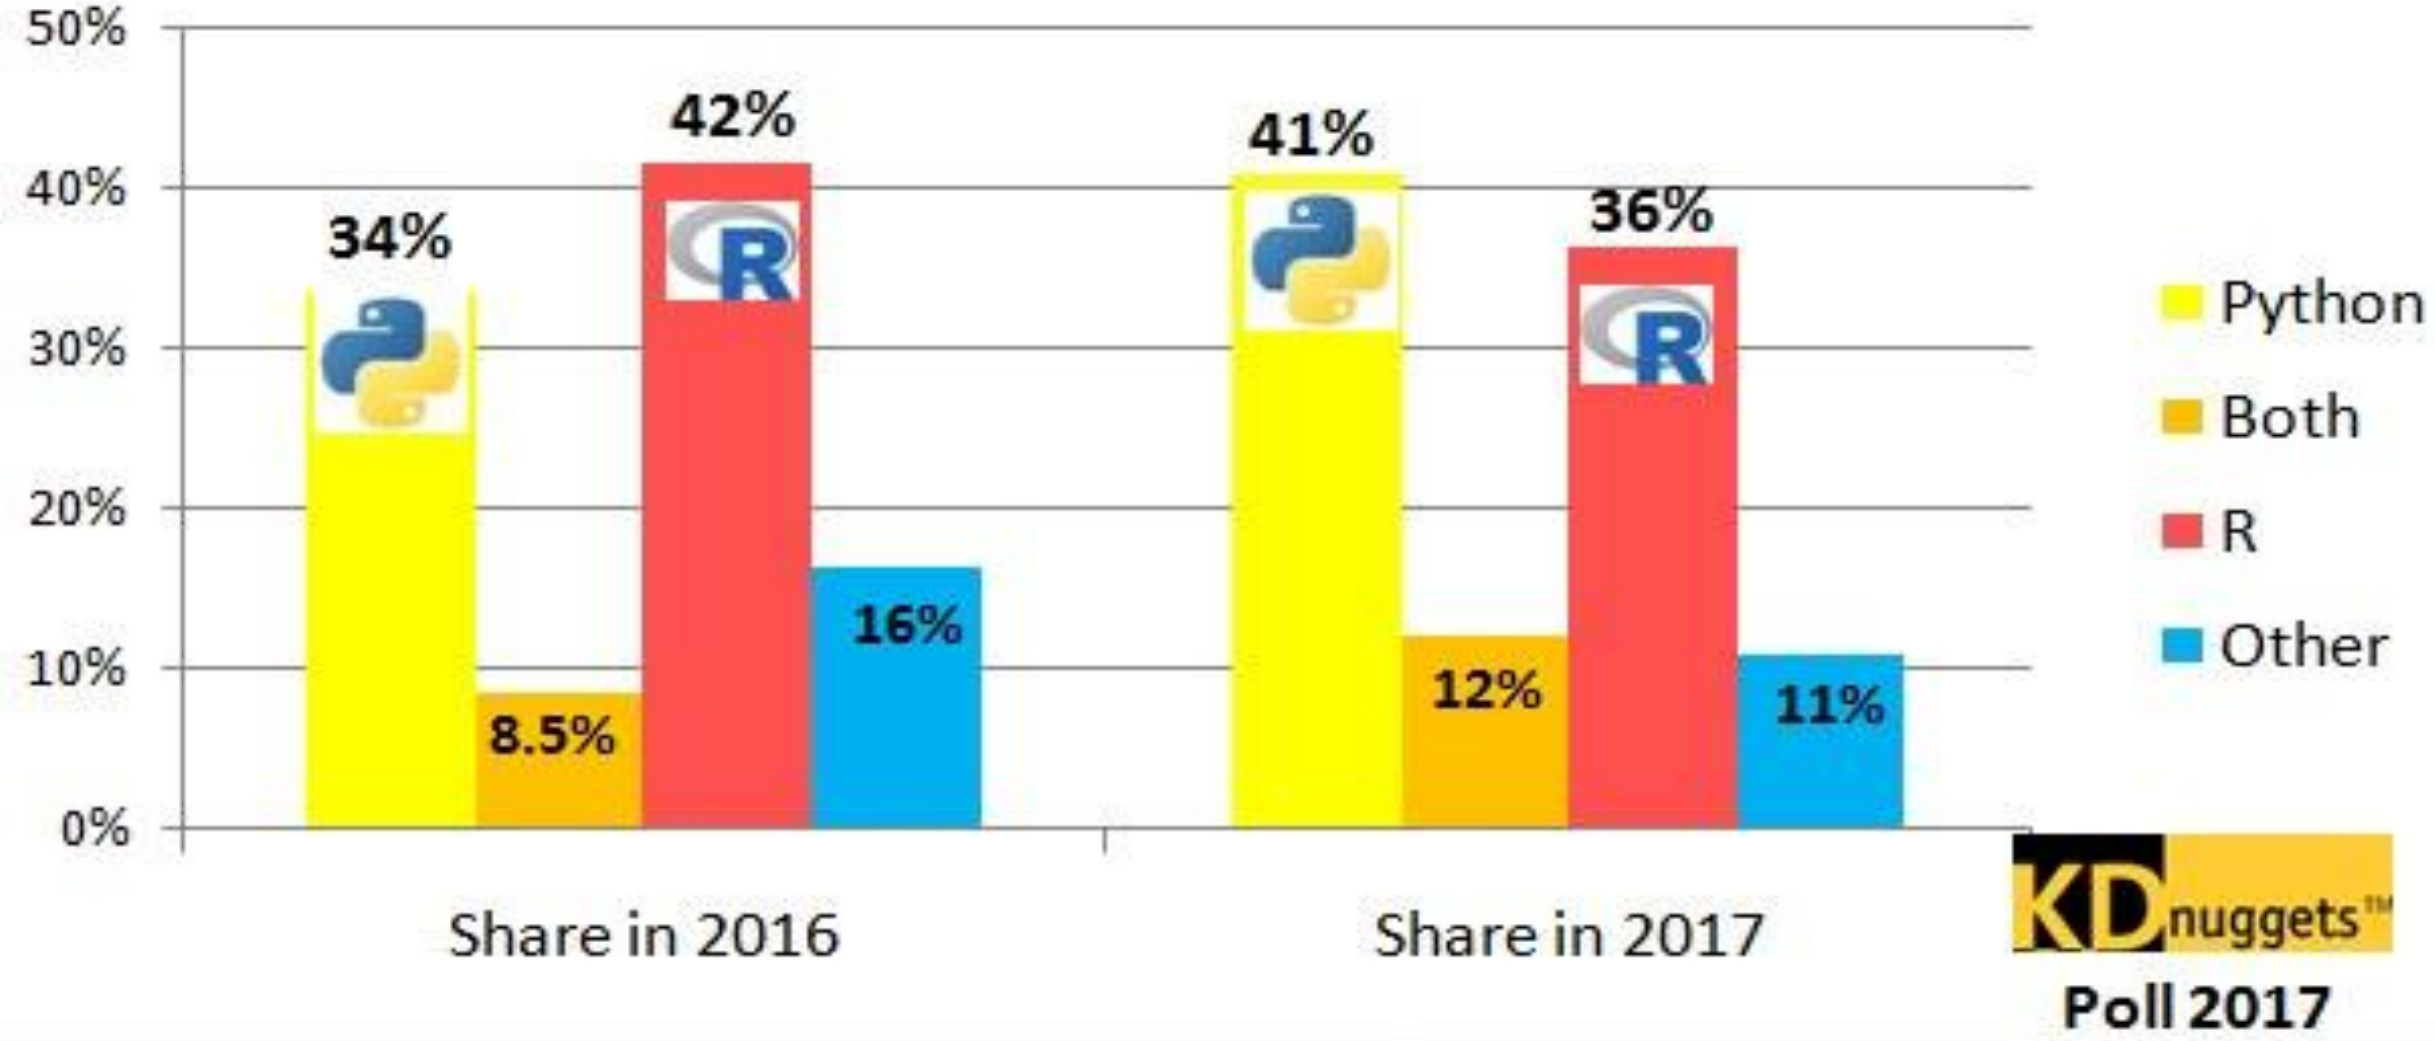

KDnuggets™

Poll 2017

Why Python for Data Science?

- Python's inherent readability and simplicity make it relatively **easy to pick up**
- Large **amount** of dedicated analytical **libraries** and **open-source communities**
- **Millions** of users who are happy to **offer advice** or suggestions
- Supports object-oriented programming, structured programming, and functional programming patterns

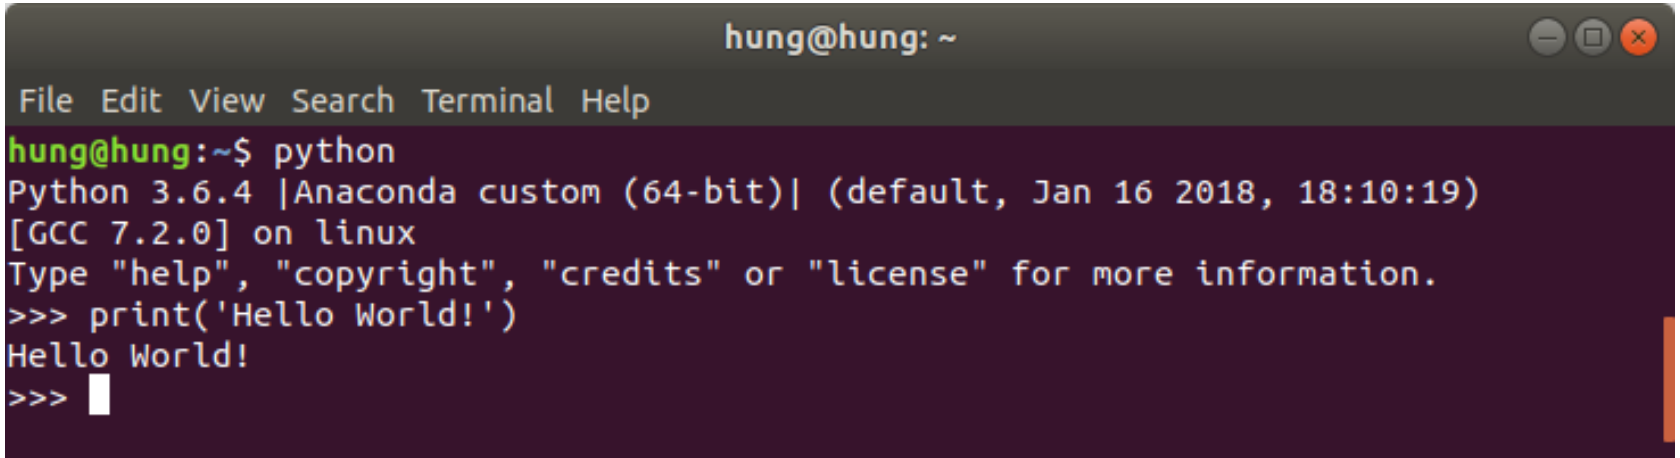

Python basic libraries for Data Science

```
Install:    $ pip install <library_name>
Use:        >>> import    <library_name>
```

IP[y]: IPython
Interactive Computing

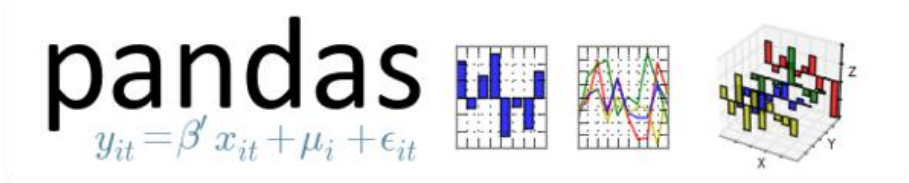

jupyter spectrogram (autosaved)

File Edit View Insert Cell Kernel Help Python 3

Markdown CellToolbar

IDE: Python Notebook

```
Install:   $ pip install jupyter
Run:       $ jupyter notebook
```

Simple spectral analysis

An illustration of the Discrete Fourier Transform

$$X_k = \sum_{n=0}^{N-1} x_n exp^{\frac{-2\pi i}{N} kn} \quad k = 0, \ldots, N-1$$

In [2]:
```
from scipy.io import wavfile
rate, x = wavfile.read('test_mono.wav')
```

And we can easily view it's spectral structure using matplotlib's builtin specgram routine:

In [5]:
```
fig, (ax1, ax2) = plt.subplots(1,2,figsize(16,5))
ax1.plot(x); ax1.set_title('Raw audio signal')
ax2.specgram(x); ax2.set_title('Spectrogram');
```

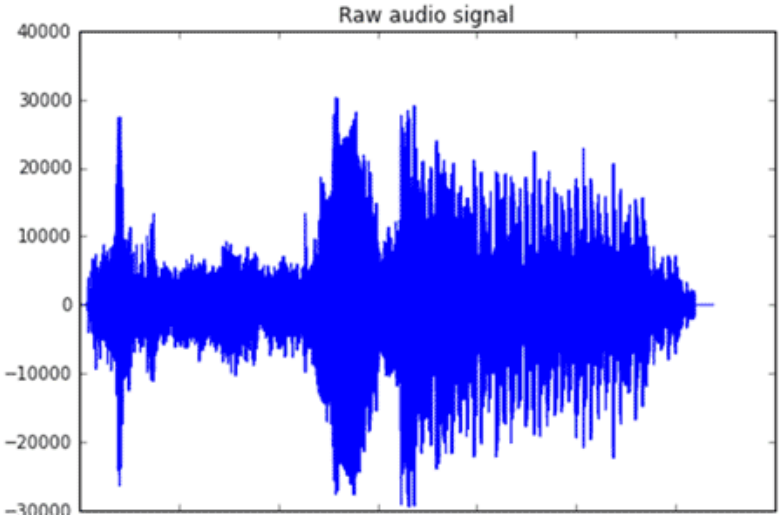

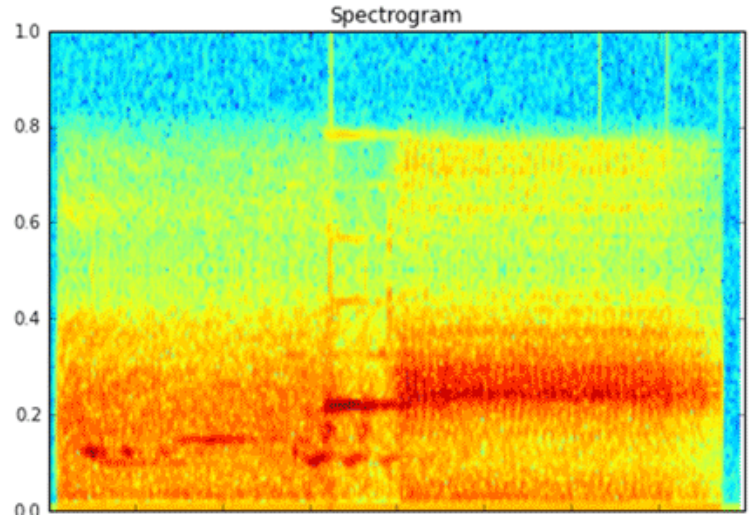

Libraries/APIs for Big Data & Deep Learning

THANK YOU!

Vietnam:	**84-2839-951-059**
North America:	**+1 844 224 4188**
Australia:	**+61 414 734 277**
Japan:	**+81 364 324 994**

Website: **http://tma-innovation.center**

Email: **innovation@tma.com.vn**